

```
In [1]: import mlflow
import numpy as np

from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout
import matplotlib.pyplot as plt
from plot_keras_history import show_history

from tensorflow.keras import regularizers

from tensorflow.keras.callbacks import ModelCheckpoint
import h5py
```


```
In [2]: (X_train, Y_train), (X_test, Y_test) = keras.datasets.mnist.load_data()
num_classes = 10
x_train = X_train.reshape(60000, 784)
x_test = X_test.reshape(10000, 784)
x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255
print(x_train.shape, 'train input samples')
print(x_test.shape, 'test input samples')
```

```
(60000, 784) train input samples
(10000, 784) test input samples
```

```
In [3]: y_train = keras.utils.to_categorical(Y_train, num_classes)
y_test = keras.utils.to_categorical(Y_test, num_classes)
print(y_train.shape, 'train output samples')
print(y_test.shape, 'test output samples')
```

```
(60000, 10) train output samples
(10000, 10) test output samples
```

```
In [4]: plt.subplot(221)
plt.imshow(X_train[np.random.randint(0, len(X_train))], cmap=plt.get_cmap(
'gray'))
plt.subplot(222)
plt.imshow(X_train[np.random.randint(0, len(X_train))], cmap=plt.get_cmap(
'gray'))
plt.subplot(223)
plt.imshow(X_train[np.random.randint(0, len(X_train))], cmap=plt.get_cmap(
'gray'))
plt.subplot(224)
plt.imshow(X_train[np.random.randint(0, len(X_train))], cmap=plt.get_cmap(
'gray'))
plt.show()
```



```
In [5]: def build_basic_nn(params):
mlp = Sequential([
    Dense(params["layer1_size"], activation=params['activation'], kernel_regularizer=params['regularizers'], input_shape=(784,)),
    Dropout(params["dropout_rate_l1"]),
    Dense(params["layer2_size"], activation=params['activation'], kernel_regularizer=params['regularizers']),
    Dropout(params["dropout_rate_l2"]),
    Dense(params["output_size"], activation='softmax')
])
return mlp
```

```
In [6]: def train_mlp(mlp, train_params, x_train, y_train):

    if train_params["use_optimiser"] == True:
        opt_new = keras.optimizers.SGD(learning_rate=train_params["learning_rate"], momentum=train_params["momentum"])
        mlp.compile(loss='categorical_crossentropy', metrics=['accuracy'], optimizer=opt_new)
    else:
        mlp.compile(loss='categorical_crossentropy', metrics=['accuracy'])

    if train_params["early_stopping"] == True:
        checkpoint = ModelCheckpoint(r"mnist-epoch-{epoch:02d}.keras")
        history = mlp.fit(x_train, y_train, epochs=train_params["num_epochs"], validation_data=(x_test, y_test), callbacks=[checkpoint])
        es = keras.callbacks.EarlyStopping(monitor='val_accuracy', min_delta=0.01, patience=2)
        mlp.fit(x_train, y_train, epochs=train_params["num_epochs"], validation_data=(x_test, y_test), callbacks=[es], batch_size=train_params["batch_size"])
    else:
        history = mlp.fit(x_train, y_train, epochs=train_params["num_epochs"], validation_data=(x_test, y_test), batch_size=train_params["batch_size"])

    show_history(history)

    return mlp
```

```
In [7]: def mlp_run(
    mlp_params,
    train_params,
    x_train,
    y_train,
    x_test,
    y_test,
):
    model = build_basic_nn(mlp_params)
    model = train_mlp(model, train_params, x_train, y_train)

    loss, acc = model.evaluate(x_train, y_train, verbose=2)
    print("Train accuracy: {:.2f}%".format(100*acc))

    loss, acc = model.evaluate(x_test, y_test, verbose=2)
    print("Test accuracy: {:.2f}%".format(100*acc))

    return model
```

```
In [8]: def predict(img, model):
    img = img.reshape(1, 784)
    img = img.astype('float32')
    img = img / 255
    probs = model.predict(img, verbose=True)
    print("Predicted Digit:", np.argmax(probs))
    return np.argmax(probs)
```


Bigger Neural Network

```
In [9]: mlp_params = {
        "layer1_size": 256,
        "dropout_rate_l1": 0,
        "layer2_size": 128,
        "dropout_rate_l2": 0,
        "output_size": 10,
        "activation": 'sigmoid',
        "regularizers": None,
    }

    train_params = dict(
        use_optimiser=False,
        learning_rate=0.1,
        momentum=0.0,
        num_epochs=10,
        early_stopping=False,
        batch_size=32,
    )

    model = mlp_run(
        mlp_params,
        train_params,
        x_train,
        y_train,
        x_test,
        y_test,
    )
```

Epoch 1/10

```
c:\Users\18arj\AppData\Local\Programs\Python\Python312\Lib\site-packages\keras\src\layers\core\dense.py:86: UserWarning: Do not pass an `input_shape`/  
`input_dim` argument to a layer. When using Sequential models, prefer using  
an `Input(shape)` object as the first layer in the model instead.
```

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

1875/1875 ————— 7s 3ms/step - accuracy: 0.8071 - loss: 0.7001 - val_accuracy: 0.9381 - val_loss: 0.2036

Epoch 2/10

1875/1875 ————— 7s 4ms/step - accuracy: 0.9424 - loss: 0.1913 - val_accuracy: 0.9533 - val_loss: 0.1465

Epoch 3/10

1875/1875 ————— 8s 4ms/step - accuracy: 0.9620 - loss: 0.1251 - val_accuracy: 0.9674 - val_loss: 0.1070

Epoch 4/10

1875/1875 ————— 6s 3ms/step - accuracy: 0.9723 - loss: 0.0919 - val_accuracy: 0.9717 - val_loss: 0.0912

Epoch 5/10

1875/1875 ————— 6s 3ms/step - accuracy: 0.9785 - loss: 0.0712 - val_accuracy: 0.9752 - val_loss: 0.0834

Epoch 6/10

1875/1875 ————— 6s 3ms/step - accuracy: 0.9831 - loss: 0.0545 - val_accuracy: 0.9780 - val_loss: 0.0706

Epoch 7/10

1875/1875 ————— 6s 3ms/step - accuracy: 0.9858 - loss: 0.0452 - val_accuracy: 0.9779 - val_loss: 0.0693

Epoch 8/10


1875/1875 ————— 6s 3ms/step - accuracy: 0.9885 - loss: 0.0374 - val_accuracy: 0.9806 - val_loss: 0.0669

Epoch 9/10

1875/1875 ————— 6s 3ms/step - accuracy: 0.9892 - loss: 0.0355 - val_accuracy: 0.9774 - val_loss: 0.0767

Epoch 10/10

1875/1875 ————— 6s 3ms/step - accuracy: 0.9914 - loss: 0.0280 - val_accuracy: 0.9819 - val_loss: 0.0665



1875/1875 - 2s - 1ms/step - accuracy: 0.9949 - loss: 0.0189

Train accuracy: 99.49%

313/313 - 0s - 1ms/step - accuracy: 0.9819 - loss: 0.0665


Test accuracy: 98.19%

```
In [10]: model.save('model_to_use.h5')
         model.save_weights('model_weights.weights.h5')

         # Save model architecture
         with open('model_to_use.json', 'w') as f:
             f.write(model.to_json())
```

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

```
In [11]: image = X_test[np.random.randint(0, len(X_test))]
         predict(image, model)
```

1/1  0s 127ms/step
Predicted Digit: 2

Out[11]: 2
